

# Week 04: Dynamic Data Structures

## Memory


1/68

Reminder:

Computer memory ... large array of consecutive data cells or bytes

- char ... 1 byte
- int, float ... 4 bytes
- double ... 8 bytes
- any\_type \* ... 4 bytes (on CSE lab computers)

Memory addresses shown in Hexadecimal notation



## C execution: Memory


2/68

An executing C program partitions memory into:

- *code* ... fixed-size, read-only region
  - contains the machine code instructions for the program
- *global data* .. fixed-size, read-write region
  - contain global variables and constant strings
- *heap* ... very large, read-write region
  - contains dynamic data structures created by malloc() (see later)
- *stack* ... dynamically-allocated data (function local vars)
  - consists of frames, one for each currently active function
  - each frame contains local variables and house-keeping info

### ... C execution: Memory

3/68



### Exercise #1: Memory Regions

4/68

```
int numbers[] = { 40, 20, 30 };
```

```
void insertionSort(int array[], int n) {
    int i, j;
    for (i = 1; i < n; i++) {
        int element = array[i];
        while (j >= 0 && array[j] > element) {
            array[j+1] = array[j];
            j--;
        }
        array[j+1] = element;
    }
}

int main(void) {
    insertionSort(numbers, 3);
    return 0;
}
```

Which memory region are the following objects located in?

1. insertionSort()
2. numbers[0]
3. n
4. array[0]
5. element

1. code
2. global
3. stack
4. global
5. stack

## Dynamic Data Structures

### Dynamic Memory Allocation

7/68

So far, we have considered *static* memory allocation

- all objects completely defined at compile-time
- sizes of all objects are known to compiler

Examples:

```
int x; // 4 bytes containing a 32-bit integer value
char *cp; // 4 bytes (on CSE machines)
// containing address of a char
typedef struct {float x; float y;} Point;
Point p; // 8 bytes containing two 32-bit float values
char s[20]; // array containing space for 20 1-byte chars
```

## ... Dynamic Memory Allocation

In many applications, fixed-size data is ok.

In many other applications, we need flexibility.

Examples:

```
char name[MAXNAME];    // how long is a name?
char item[MAXITEMS];   // how high can the stack grow?
char dictionary[MAXWORDS][MAXWORDLENGTH];
                    // how many words are there?
                    // how long is each word?
```

With fixed-size data, we need to guess sizes ("large enough").

## ... Dynamic Memory Allocation

Fixed-size memory allocation:

- allocate as much space as we might ever possibly need

Dynamic memory allocation:

- allocate as much space as we actually need
- determine size based on inputs

But how to do this in C?

- all data allocation methods so far are "static"
  - however, stack data (when calling a function) is created dynamically (size is known)

## Dynamic Data Example

Problem:

- read integer data from standard input (keyboard)
- first number tells how many numbers follow
- rest of numbers are read into a vector
- subsequent computation uses vector (e.g. sorts it)

Example input: 6 25 -1 999 42 -16 64

How to define the vector?

## ... Dynamic Data Example

Suggestion #1: allocate a large vector; use only part of it

```
#define MAXElems 1000
```

8/68

```
// how many elements in the vector
int numberOfElems;
scanf("%d", &numberOfElems);
assert(numberOfElems <= MAXElems);

// declare vector and fill with user input
int i, vector[MAXElems];
for (i = 0; i < numberOfElems; i++)
    scanf("%d", &vector[i]);
```

Works ok, unless too many numbers; usually wastes space.

Recall that `assert()` terminates program with standard error message if test fails.

## ... Dynamic Data Example

Suggestion #2: create vector after count read in

```
#include <stdlib.h>
```

```
// how many elements in the vector
int numberOfElems;
scanf("%d", &numberOfElems);

// declare vector and fill with user input
int i, *vector;
size_t numberOfBytes;
numberOfBytes = numberOfElems * sizeof(int);
```

```
vector = malloc(numberOfBytes);
assert(vector != NULL);
```

```
for (i = 0; i < numberOfElems; i++)
    scanf("%d", &vector[i]);
```

Works unless the *heap* is already full (very unlikely)

Reminder: because of pointer/array connection `&vector[i] == vector+i`


## The `malloc()` function

Recall memory usage within C programs:

11/68

12/68

13/68



- its return value is a *generic* pointer (`void *`)
- the return value must *always* be checked (may be NULL)

Required size is determined by `#Elements * sizeof(ElementType)`

## Exercise #2: Dynamic Memory Allocation

17/68

Write code to

1. create space for 1,000 speeding tickets (cf. Lecture Week 1)
2. create a dynamic  $m \times n$ -matrix of floating point numbers, given  $m$  and  $n$

How many bytes need to be reserved in each case?

### ... The `malloc()` function

14/68

`malloc()` function interface

```
void *malloc(size_t n);
```

What the function does:

- attempts to reserve a block of  $n$  bytes in the *heap*
- returns the address of the start of this block
- if insufficient space left in the heap, returns NULL

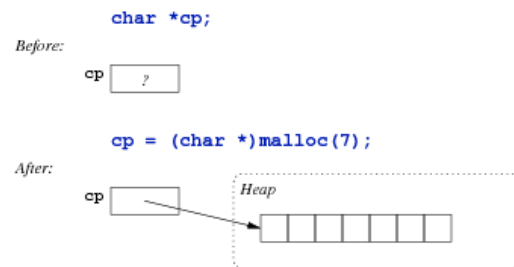
Note: `size_t` is essentially an unsigned `int`

- but has specialised interpretation of applying to memory sizes measured in bytes

### ... The `malloc()` function

15/68

Example use of `malloc`:



### ... The `malloc()` function

16/68

Things to note about `void *malloc(size_t):`

- it is defined as part of `stdlib.h`
- its parameter is a size in units of *bytes*

1. Speeding tickets:

```
typedef struct {
    int day, month, year; } DateT;
typedef struct {
    int hour, minute; } TimeT;
typedef struct {
    char plate[7]; DateT d; TimeT t; } TicketT;
```

```
TicketT *tickets = malloc(1000 * sizeof(TicketT));
assert(tickets != NULL);
```

28,000 bytes allocated

2. Matrix:

```
float **matrix = malloc(m * sizeof(float *));
assert(matrix != NULL);
int i;
for (i = 0; i < m; i++) {
    matrix[i] = malloc(n * sizeof(float));
    assert(matrix[i] != NULL);
}
```

$4m + 4mn$  bytes allocated

## Exercise #3: Memory Regions

19/68

Which memory region is `tickets` located in? What about `*tickets`?

1. `tickets` is a variable located in the stack
2. `*tickets` is in the heap (after `malloc`'ing memory)

## ... The `malloc()` function

21/68

`malloc()` returns a pointer to a data object of some kind.

Things to note about objects allocated by `malloc()`:

- they exist until explicitly removed (program-controlled lifetime)
- they are *accessible* while some variable references them
- if no active variable references an object, it is *garbage*

The function `free()` releases objects allocated by `malloc()`

## ... The `malloc()` function

22/68

Usage of `malloc()` should always be guarded:

```
int *vector, length, i;
...
vector = malloc(length*sizeof(int));
// but malloc() might fail to allocate
assert(vector != NULL);
// now we know it's safe to use vector[]
for (i = 0; i < length; i++) {
    ... vector[i] ...
}
```

Alternatively:

```
int *vector, length, i;
...
vector = malloc(length*sizeof(int));
// but malloc() might fail to allocate
if (vector == NULL) {
    fprintf(stderr, "Out of memory\n");
    exit(1);
}
// now we know its safe to use vector[]
for (i = 0; i < length; i++) {
    ... vector[i] ...
}
```

- `fprintf(stderr, ...)` outputs text to a stream called **`stderr`** (the screen, by default)
- `exit(v)` terminates the program with return value `v`

## Memory Management

23/68

### **`void free(void *ptr)`**

- releases a block of memory allocated by `malloc()`
- `*ptr` is a dynamically allocated object
- if `*ptr` was not `malloc()`'d, chaos will follow

Things to note:

- the contents of the memory block are not changed
- all pointers to the block still exist, but are not valid
- the memory may be re-used as soon as it is `free()`'d

## ... Memory Management

24/68

### **Warning! Warning! Warning! Warning!**

Careless use of `malloc()` / `free()` / pointers

- can mess up the data in the heap
- so that later `malloc()` or `free()` cause run-time errors
- possibly well after the original error occurred

Such errors are **very difficult** to track down and debug.

Must be **very careful** with your use of `malloc()` / `free()` / pointers.

## ... Memory Management

25/68

If an uninitialised or otherwise invalid pointer is used, or an array is accessed with a negative or out-of-bounds index, one of a number of things might happen:

- program aborts immediately with a "segmentation fault"
- a mysterious failure much later in the execution of the program
- incorrect results, but no obvious failure
- correct results, but maybe not always, and maybe not when executed on another day, or another machine

The first is the most desirable, but cannot be relied on.

## ... Memory Management

26/68

Given a pointer variable:

- you can check whether its value is `NULL`
- you can (maybe) check that it is an address
- you **cannot** check whether it is a valid address

## ... Memory Management

27/68

Typical usage pattern for dynamically allocated objects:

```
// single dynamic object e.g. struct
Type *ptr = malloc(sizeof(Type));
assert(ptr != NULL);
... use object referenced by ptr e.g. ptr->name ...
```

```
free(ptr);

// dynamic array with "nelems" elements
int nelems = NumberOfElements;
ElemType *arr = malloc(nelems*sizeof(ElemType));
assert(arr != NULL);
... use array referenced by arr e.g. arr[4] ...
free(arr);
```

---

## Memory Leaks

28/68

Well-behaved programs do the following:

- allocate a new object via `malloc()`
- use the object for as long as needed
- `free()` the object when no longer needed

A program which does not `free()` each object before the last reference to it is lost contains a *memory leak*.

Such programs may eventually exhaust available heap space.

---

### Exercise #4: Dynamic Arrays

29/68

Write a C-program that

- prompts the user to input a positive number  $n$
- allocates memory for two  $n$ -dimensional floating point vectors **a** and **b**
- prompts the user to input  $2n$  numbers to initialise these vectors
- computes and outputs the inner product of **a** and **b**
- frees the allocated memory

---

## Sidetrack: Standard I/O Streams, Redirects

30/68

Standard file streams:

- **stdin** ... standard input, by default: keyboard
- **stdout** ... standard output, by default: screen
- **stderr** ... standard error, by default: screen

- `fprintf(stdout, ...)` has the same effect as `printf(...)`
- `fprintf(stderr, ...)` often used to print error messages

Executing a C program causes `main(...)` to be invoked

- with `stdin`, `stdout`, `stderr` already open for use

---

### ... Sidetrack: Standard I/O Streams, Redirects

31/68

The streams `stdin`, `stdout`, `stderr` can be *redirected*

- redirecting `stdin`

```
prompt$ myprog < input.data
```

- redirecting `stdout`

```
prompt$ myprog > output.data
```

- redirecting `stderr`

```
prompt$ myprog 2> error.data
```

---

## Abstract Data Structures: ADTs

### Abstract Data Types

33/68

Reminder: An *abstract data type* is ...

- an approach to implementing data types
- separates *interface* from *implementation*
- users of the ADT see only the interface
- builders of the ADT provide an implementation

E.g. does a client want/need to know how a Stack is implemented?

- ADO = *abstract data object* (e.g. a single stack)
- ADT = *abstract data type* (e.g. stack data type)

---

### ... Abstract Data Types

34/68

ADT *interface* provides

- an *opaque* user-view of the data structure (e.g. `stack *`)
- function signatures (prototypes) for all operations
- semantics of operations (via documentation)
- a contract between ADT and its clients

ADT *implementation* gives

- concrete definition of the data structure
- function implementations for all operations
- ... including for *creation* and *destruction* of instances of the data structure

ADTs are important because ...

- facilitate decomposition of complex programs
- make implementation changes invisible to clients
- improve readability and structuring of software

# Stack as ADT

35/68

Interface (in `stack.h`)

```
// provides an opaque view of ADT
typedef struct StackRep *stack;

// set up empty stack
stack newStack();
// remove unwanted stack
void dropStack(stack);
// check whether stack is empty
int StackIsEmpty(stack);
// insert an int on top of stack
void StackPush(stack, int);
// remove int from top of stack
int StackPop(stack);
```

ADT *stack* defined as a *pointer* to an *unspecified* struct named `StackRep`

# Sidetrack: Defining Structures

36/68

Structures can be defined in two different styles:

```
typedef struct { int day, month, year; } DateT;
// which would be used as
DateT somedate;
```

// or

```
struct date { int day, month, year; };
// which would be used as
struct date anotherdate;
```

The definitions produce objects with identical structures.

It is possible to combine both styles:

```
typedef struct date { int day, month, year; } DateT;
// which could be used as
DateT      date1, *dateptr1;
struct date date2, *dateptr2;
```

# Static/Dynamic Sequences

37/68

Previously we have used an *array* to implement a stack

- fixed size collection of heterogeneous elements
- can be accessed via index or via "moving" pointer

The "fixed size" aspect is a potential problem:

- how big to make the (dynamic) array? (big ... just in case)
- what to do if it fills up?


The rigid sequence is another problems:

- inserting/deleting an item in middle of array

# ... Static/Dynamic Sequences

38/68


Inserting a value into a sorted array (`insert(a, &n, 4)`):



# ... Static/Dynamic Sequences

39/68

Deleting a value from a sorted array (`delete(a, &n, 3)`):




# Dynamic Sequences

40/68

The problems with using arrays can be solved by

- allocating elements individually
- linking them together as a "chain"



Benefits:

- insertion/deletion have minimal effect on list overall
- only use as much space as needed for values

# Self-referential Structures

41/68

To realise a "chain of elements", need a *node* containing

- a value
- a link to the next node

In C, we can define such nodes as:

```
typedef struct node {
    int data;
    struct node *next;
} NodeT;
```

## ... Self-referential Structures

42/68

Note that the following definition does not work:

```
typedef struct {
    int data;
    NodeT *next;
} NodeT;
```

Because NodeT is not yet known (to the compiler) when we try to use it to define the type of the next field.

The following is also illegal in C:

```
struct node {
    int data;
    struct node recursive;
};
```

Because the size of the structure would have to satisfy `sizeof(struct node) = sizeof(int) + sizeof(struct node)` = ∞.


# Linked Lists in C

## Linked Lists

44/68

To represent a chained (linked) list of nodes:

- we need a pointer to the first node
- each node contains a pointer to the next node
- the next pointer in the last node is NULL



## ... Linked Lists

45/68

Linked lists are more flexible than arrays:

- values do not have to be adjacent in memory
- values can be rearranged simply by altering pointers
- the number of values can change dynamically
- values can be added or removed in any order

Disadvantages:

- it is not difficult to get pointer manipulations wrong
- each value also requires storage for next pointer

# Memory Storage for Linked Lists

46/68

Linked list nodes are typically located in the heap

- because nodes are dynamically created

Variables containing pointers to list nodes

- are likely to be local variables (in the stack)

Pointers to the start of lists are often

- passed as parameters to function
- returned as function results

## ... Memory Storage for Linked Lists

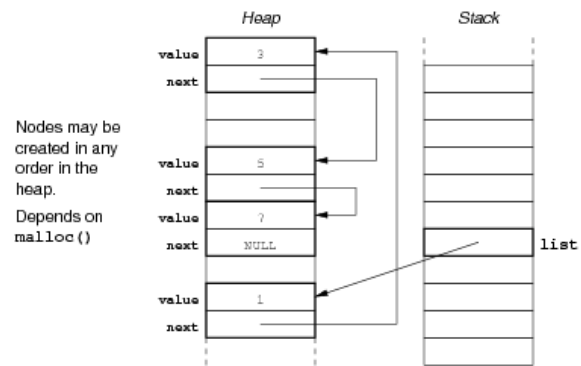
47/68

Create a new list node:

```
NodeT *makeNode(int v) {
    NodeT *new = malloc(sizeof(NodeT));
    assert(new != NULL);
    new->data = v;           // initialise data
    new->next = NULL;        // initialise link to next node
    return new;              // return pointer to new node
}
```

## ... Memory Storage for Linked Lists

48/68



## Exercise #5: Creating a linked list

49/68

Write C-code to create a linked list of three nodes with values 1, 42 and 9024.

```
NodeT *list = makeNode(1);
list->next = makeNode(42);
list->next->next = makeNode(9024);
```

## Iteration over Linked Lists

51/68

When manipulating list elements

- typically have pointer `p` to current node (`NodeT *p`)
- to access the data in current node: `p->data`
- to get pointer to next node: `p->next`

To iterate over a linked list:

- set `p` to point at first node (head)
- examine node pointed to by `p`
- change `p` to point to next node
- stop when `p` reaches end of list (NULL)

## ... Iteration over Linked Lists

52/68

Standard method for scanning all elements in a linked list:

```
NodeT *list; // pointer to first Node in list
NodeT *p;    // pointer to "current" Node in list
```

```
p = list;
while (p != NULL) {
    ... do something with p->data ...
}
```


```
p = p->next;
}

// which is frequently written as

for (p = list; p != NULL; p = p->next) {
    ... do something with p->data ...
}
```


## ... Iteration over Linked Lists

53/68



## ... Iteration over Linked Lists

54/68



## ... Iteration over Linked Lists

55/68

Check if list contains an element:

```
int inLL(NodeT *list, int d) {
```



```

NodeT *p;
for (p = list; p != NULL; p = p->next)
    if (p->data == d)        // element found
        return 1;
return 0;                    // element not in list
}

```

Print all elements:

```

void showLL(NodeT *list) {
    NodeT *p;
    for (p = list; p != NULL; p = p->next)
        printf("%6d", p->data);
}

```

## Exercise #6: Traversing a linked list

56/68

What does this code do?

```

1  NodeT *p = list;
2  while (p != NULL) {
3      printf("%6d", p->data);
4      if (p->next != NULL)
5          p = p->next->next;
6      else
7          p = NULL;
8  }

```

What is the purpose of the conditional statement in line 4?

Every second list element is printed.

If `*p` happens to be the last element in the list, then `p->next->next` does not exist. The if-statement ensures that we do not attempt to assign an invalid address to `p` in line 5.

## Exercise #7: Traversing a linked list

58/68

Rewrite `showLL()` as a recursive function.

```

void printLL(NodeT *list) {
    if (list != NULL) {
        printf("%6d", list->data);
        printLL(list->next);
    }
}

```

## Modifying a Linked List

60/68

Insert a new element at the beginning:

```

NodeT *insertLL(NodeT *list, int d) {
    NodeT *new = makeNode(d); // create new list element
    new->next = list;          // link to beginning of list
    return new;                // new element is new head
}

```

Delete the first element:

```

NodeT *deleteHead(NodeT *list) {
    assert(list != NULL); // ensure list is not empty
    NodeT *head = list;   // remember address of first element
    list = list->next;     // move to second element
    free(head);           // return pointer to second element
    return list;
}

```

What would happen if we didn't free the memory pointed to by `head`?

## ... Modifying a Linked List

61/68

Delete a specific element (recursive version):

```

NodeT *deleteLL(NodeT *list, int d) {
    if (list == NULL) { // element not in list
        return list;

        } else if (list->data == d) {
            return deleteHead(list); // delete first element

        } else { // delete element in tail list
            list->next = deleteLL(list->next, d);
            return list;
        }
}

```

## Exercise #8: Freeing a list

62/68

Write a C-function to destroy an entire list.

Iterative version:

```

void freeLL(NodeT *list) {
    NodeT *p;

    p = list;
    while (p != NULL) {
        NodeT *temp = p->next;
        free(p);
        p = temp;
    }
}

```

```
}  
}
```

Why do we need the extra variable temp?

## Stack ADT Implementation

64/68

Linked list implementation (stack.c):

Remember: stack.h includes `typedef struct StackRep *stack;`

```
#include <stdlib.h>  
#include <assert.h>  
#include "stack.h"  
  
typedef struct node {  
    int data;  
    struct node *next;  
} NodeT;  
  
typedef struct StackRep {  
    int height; // #elements on stack  
    NodeT *top; // ptr to first element  
} StackRep;  
  
// set up empty stack  
stack newStack() {  
    stack S = malloc(sizeof(StackRep));  
    S->height = 0;  
    S->top = NULL;  
    return S;  
}  
  
// remove unwanted stack  
void dropStack(stack S) {  
    NodeT *curr = S->top;  
    while (curr != NULL) { // free the list  
        NodeT *temp = curr->next;  
        free(curr);  
        curr = temp;  
    }  
    free(S); // free the stack rep  
}  
  
// check whether stack is empty  
int StackIsEmpty(stack S) {  
    return (S->height == 0);  
}  
  
// insert an int on top of stack  
void StackPush(stack S, int v) {  
    NodeT *new = malloc(sizeof(NodeT));  
    assert(new != NULL);  
    new->data = v;  
    // insert new element at top  
    new->next = S->top;  
    S->top = new;  
    S->height++;  
}  
  
// remove int from top of stack  
int StackPop(stack S) {  
    assert(S->height > 0);  
    NodeT *head = S->top;  
    // second list element becomes new top  
    S->top = S->top->next;  
    S->height--;  
    // read data off first element, then free  
    int d = head->data;  
    free(head);  
    return d;  
}
```

## Summary: Memory Management Functions

65/68

`void *malloc(size_t nbytes)`

- aim: allocate some memory for a data object
- attempt to allocate a block of memory of size `nbytes` in the heap
- if successful, returns a pointer to the start of the block
- if insufficient space in heap, returns `NULL`

Things to note:

- the location of the memory block within heap is random
- the initial contents of the memory block are random

## ... Summary: Memory Management Functions

`void free(void *ptr)`

- releases a block of memory allocated by `malloc()`
- `*ptr` is the start of a dynamically allocated object
- if `*ptr` was not `malloc()`'d, chaos will ensue

Things to note:

- the contents of the memory block are not changed
- all pointers to the block still exist, but are not valid
- the memory may be re-used as soon as it is `free()`'d

## Tips for Week 4 Problem Set

67/68

Main theme: *Dynamic data structures*

- Test your understanding of memory allocation and deallocation
- Create, use and free a dynamic array (Exercise 3)
- Think about how to design and use a dynamic queue ADT (Exercise 4)
- Design and implement functions for dynamic linked lists (Exercise 5)

```
prompt$ ./llbuild  
Enter an integer: 12  
Enter an integer: 34  
Enter an integer: 56  
Enter an integer: quit  
Finished. List is 12->34->56
```

- Challenge Exercise: wrack your brain — split linked list in two halves without traversing it twice

## Summary

68/68

- Memory management
- Dynamic data structures
- Linked lists

- Suggested reading:
  - Moffat, Ch.10.1-10.2
  - Sedgewick, Ch.3.3-3.5,4.4,4.6

Produced: 13 Aug 2018

66/68